

LSTM

CLASS torch.nn.LSTM(*input_size*, *hidden_size*, *num_layers*=1, *bias*=True, *batch_first*=False, *dropout*=0.0, *bidirectional*=False, *proj_size*=0, *device*=None, *dtype*=None) [SOURCE]

Apply a multi-layer long short-term memory (LSTM) RNN to an input sequence. For each element in the input sequence, each layer computes the following function:

$$\begin{aligned} i_t &= \sigma(W_{ii}x_t + b_{ii} + W_{hi}h_{t-1} + b_{hi}) \\ f_t &= \sigma(W_{if}x_t + b_{if} + W_{hf}h_{t-1} + b_{hf}) \\ g_t &= \tanh(W_{ig}x_t + b_{ig} + W_{hg}h_{t-1} + b_{hg}) \\ o_t &= \sigma(W_{io}x_t + b_{io} + W_{ho}h_{t-1} + b_{ho}) \\ c_t &= f_t \odot c_{t-1} + i_t \odot g_t \\ h_t &= o_t \odot \tanh(c_t) \end{aligned}$$

$$it = \sigma(W_{iixt} + b_{iix} + W_{hiht-1} + b_{hih})ft = \sigma(W_{ifxt} + b_{ifx} + W_{hfht-1} + b_{hfh})gt = \tanh(W_{igxt} + b_{igx} + W_{hght-1} + b_{hgh})ot = \sigma(W_{ioxt} + b_{iox} + W_{hoht-1} + b_{hoh})ct = ft \odot ct-1 + it$$

$$\odot ght = ot \odot \tanh(ct)$$

where h_t , ht is the hidden state at time t , c_t , ct is the cell state at time t , x_t , xt is the input at time t , h_{t-1} , $ht-1$ is the hidden state of the layer at time $t-1$ or the initial hidden state at time o , and i , it , f , ft , g , gt , o , ot are the input, forget, cell, and output gates, respectively. σ is the sigmoid function, and \odot is the Hadamard product.

In a multilayer LSTM, the input $x_t^{(l)}$ of the l -th layer ($l \geq 2$) is the hidden state $h_t^{(l-1)}$ of the previous layer multiplied by dropout $\delta_t^{(l-1)}$ where each $\delta_t^{(l-1)}$ is a Bernoulli random variable which is 0 with probability `dropout`.

If `proj_size > 0` is specified, LSTM with projections will be used. This changes the LSTM cell in the following way. First, the dimension of h_t will be changed from `hidden_size` to `proj_size` (dimensions of W_{hi} will be changed accordingly). Second, the output hidden state of each layer will be multiplied by a learnable projection matrix: $\tilde{h}_t = W_{hr}h_t$. Note that as a consequence of this, the output of LSTM network will be of different shape as well. See Inputs/Outputs sections below for exact dimensions of all variables. You can find more details in <https://arxiv.org/abs/1402.1128>.

Parameters

- input_size** – The number of expected features in the input x
- hidden_size** – The number of features in the hidden state h
- num_layers** – Number of recurrent layers. E.g., setting `num_layers=2` would mean stacking two LSTMs together to form a *stacked LSTM*, with the second LSTM taking in outputs of the first LSTM and computing the final results. Default: 1
- bias** – If `False`, then the layer does not use bias weights b_{ih} and b_{hh} . Default: `True`
- batch_first** – If `True`, then the input and output tensors are provided as *(batch, seq, feature)* instead of *(seq, batch, feature)*. Note that this does not apply to hidden or cell states. See the Inputs/Outputs sections below for details. Default: `False`
- dropout** – If non-zero, introduces a *Dropout* layer on the outputs of each LSTM layer except the last layer, with dropout probability equal to `dropout`. Default: 0
- bidirectional** – If `True`, becomes a bidirectional LSTM. Default: `False`
- proj_size** – If `> 0`, will use LSTM with projections of corresponding size. Default: 0

Inputs: input, (h_0, c_0)

- input**: tensor of shape (L, H_{in}) for unbatched input, (L, N, H_{in}) when `batch_first=False` or (N, L, H_{in}) when `batch_first=True` containing the features of the input sequence. The input can also be a packed variable length sequence. See `torch.nn.utils.rnn.pack_padded_sequence()` or `torch.nn.utils.rnn.pack_sequence()` for details.
- h_0**: tensor of shape $(D * num_layers, H_{out})$ for unbatched input or $(D * num_layers, N, H_{out})$ containing the initial hidden state for each element in the input sequence. Defaults to zeros if (h_0, c_0) is not provided.
- c_0**: tensor of shape $(D * num_layers, H_{cell})$ for unbatched input or $(D * num_layers, N, H_{cell})$ containing the initial cell state for each element in the input sequence. Defaults to zeros if (h_0, c_0) is not provided.

where:

$$\begin{aligned} N &= \text{batch size} \\ L &= \text{sequence length} \\ D &= 2 \text{ if bidirectional=True otherwise } 1 \\ H_{in} &= \text{input_size} \\ H_{cell} &= \text{hidden_size} \\ H_{out} &= \text{proj_size if proj_size > 0 otherwise hidden_size} \end{aligned}$$

$$N = L = D = H_{in} = H_{cell} = H_{out} = \text{batch size} \text{sequence length} 2 \text{ if bidirectional=True otherwise } 1 \text{input_size hidden_size proj_size if proj_size > 0 otherwise hidden_size}$$

Outputs: output, (h_n, c_n)

- output**: tensor of shape $(L, D * H_{out})$ for unbatched input, $(L, N, D * H_{out})$ when `batch_first=False` or $(N, L, D * H_{out})$ when `batch_first=True` containing the output features (\tilde{h}_t) from the last layer of the LSTM, for each t . If a `torch.nn.utils.rnn.PackedSequence` has been given as the input, the output will also be a packed sequence. When `bidirectional=True`, *output* will contain a concatenation of the forward and reverse hidden states at each time step in the sequence.
- h_n**: tensor of shape $(D * num_layers, H_{out})$ for unbatched input or $(D * num_layers, N, H_{out})$ containing the final hidden state for each element in the sequence. When `bidirectional=True`, h_n will contain a concatenation of the final forward and reverse hidden states, respectively.
- c_n**: tensor of shape $(D * num_layers, H_{cell})$ for unbatched input or $(D * num_layers, N, H_{cell})$ containing the final cell state for each element in the sequence. When `bidirectional=True`, c_n will contain a concatenation of the final forward and reverse cell states, respectively.

Variables

- weight_ih_l[k]** – the learnable input-hidden weights of the k^{th} layer $(W_{ij}|W_{if}|W_{ig}|W_{io})$, of shape $(4*hidden_size, input_size)$ for $k=0$. Otherwise, the shape is $(4*hidden_size, num_directions * hidden_size)$. If `proj_size > 0` was specified, the shape will be $(4*hidden_size, num_directions * proj_size)$ for $k>0$
- weight_hh_l[k]** – the learnable hidden-hidden weights of the k^{th} layer $(W_{hi}|W_{hf}|W_{hg}|W_{ho})$, of shape $(4*hidden_size, hidden_size)$. If `proj_size > 0` was specified, the shape will be $(4*hidden_size, proj_size)$.
- bias_ih_l[k]** – the learnable input-hidden bias of the k^{th} layer $(b_{ij}|b_{if}|b_{ig}|b_{io})$, of shape $(4*hidden_size)$
- bias_hh_l[k]** – the learnable hidden-hidden bias of the k^{th} layer $(b_{hi}|b_{hf}|b_{hg}|b_{ho})$, of shape $(4*hidden_size)$
- weight_hr_l[k]** – the learnable projection weights of the k^{th} layer of shape $(proj_size, hidden_size)$. Only present when `proj_size > 0` was specified.
- weight_hh_l[k]_reverse** – Analogous to *weight_ih_l[k]* for the reverse direction. Only present when `bidirectional=True`.
- weight_hh_l[k]_reverse** – Analogous to *weight_hh_l[k]* for the reverse direction. Only present when `bidirectional=True`.
- bias_ih_l[k]_reverse** – Analogous to *bias_ih_l[k]* for the reverse direction. Only present when `bidirectional=True`.
- bias_hh_l[k]_reverse** – Analogous to *bias_hh_l[k]* for the reverse direction. Only present when `bidirectional=True`.
- weight_hr_l[k]_reverse** – Analogous to *weight_hr_l[k]* for the reverse direction. Only present when `bidirectional=True` and `proj_size > 0` was specified.

* NOTE

All the weights and biases are initialized from $U(-\sqrt{k}, \sqrt{k})U(-\frac{1}{\sqrt{k}}, \frac{1}{\sqrt{k}})$ where $k = \frac{1}{max(hidden_size, proj_size)}$

* NOTE

For bidirectional LSTMs, forward and backward are directions 0 and 1 respectively. Example of splitting the output layers when `batch_first=False`:

```
output.view(seq_len, batch, num_directions, hidden_size)
```

• NOTE

For bidirectional LSTMs, `h_n` is not equivalent to the last element of `output`; the former contains the final forward and reverse hidden states, while the latter contains the final forward hidden state and the initial reverse hidden state.

• NOTE

`batch_first` argument is ignored for unbatched inputs.

• NOTE

`proj_size` should be smaller than `hidden_size`.

• WARNING

There are known non-determinism issues for RNN functions on some versions of cuDNN and CUDA. You can enforce deterministic behavior by setting the following environment variables:

On CUDA 10.1, set environment variable `CUDA_LAUNCH_BLOCKING=1`. This may affect performance.

On CUDA 10.2 or later, set environment variable (note the leading colon symbol) `CUBLAS_WORKSPACE_CONFIG=:16:8` or `CUBLAS_WORKSPACE_CONFIG=:4096:2`.

See the [cuDNN 8 Release Notes](#) for more information.

• NOTE

If the following conditions are satisfied: 1) cudnn is enabled, 2) input data is on the GPU 3) input data has dtype `torch.float16` 4) V100 GPU is used, 5) input data is not in `PackedSequence` format persistent algorithm can be selected to improve performance.

Examples:

```
>>> rnn = nn.LSTM(10, 20, 2)
>>> input = torch.randn(5, 3, 10)
>>> h0 = torch.randn(2, 3, 20)
>>> c0 = torch.randn(2, 3, 20)
>>> output, (hn, cn) = rnn(input, (h0, c0))
```

[◀ Previous](#)

[Next ▶](#)

Docs

Access comprehensive developer documentation for PyTorch
[View Docs](#)

Tutorials

Get in-depth tutorials for beginners and advanced developers
[View Tutorials](#)

Resources

Find development resources and get your questions answered
[View Resources](#)

PyTorch

Get Started

Features

Ecosystem

Blog

Contributing

Resources

Tutorials

Docs

Discuss

Github Issues

Brand Guidelines

Stay up to date

Facebook

Twitter

YouTube

LinkedIn

PyTorch Podcasts

Spotify

Apple

Google

Amazon

[Terms](#) | [Privacy](#)